The Distributed Retired Traveling Salesman Problem

Daniel Seita   DTS1@WILLIAMS.EDU
Ziang Zhang   ZZ2@WILLIAMS.EDU
Department of Computer Science, Williams College, Williamstown, MA 01267 USA

Abstract

The use of major online travel agencies has made scheduling long-term travel much easier by allowing users to easily identify a set of flights in just a few clicks. Current travel agencies allow users to plan out long-term trips involving flights to more than two cities, but they require specific arrival and departure dates for each city and a city ordering. We present a system that does not burden the user with these decisions. Specifically, our code takes in two required inputs: a list of cities the user wishes to travel to, and a date range over which they are willing to travel, and outputs the cheapest set of flights within that range that form a valid route. It essentially solves a harder version of the Traveling Salesman Problem since costs are not constant between two cities. Under several weak assumptions, and assuming that the number of cities and days is sufficiently limited, then our algorithm should successfully find the cheapest cost flight in a reasonable amount of time. We present the theoretical and systematic components of the project and discuss empirical results.

1. Introduction and Motivation

The use of major online travel agencies, such as Travelocity[1] and Kayak[2], has made scheduling long-term travel much easier by allowing users to easily select a set of flights to purchase tickets from. People use a combination of factors to help them make their decision, such as the total price of the flights and the days they wish to land and depart from a city. Current travel agencies allow users to plan out trips involving airlines to multiple destinations (see Figure 1 for

---

[1] www.travelocity.com
[2] www.kayak.com

---

an example using Travelocity), but they require (1) specific dates for each city and (2) an ordering.

This limitation can make searching for flight routes time-consuming for users who are not restricted to arriving at cities on particular dates. For instance, suppose one resides in New York City (NYC) and wants to schedule a summer trip to Paris, London, and Tokyo from June 1 to June 20, and it does not matter to him how long he stays in a city. Suppose that it also does not matter the ordering that he arrives at the cities. Thus (for now) we will only care about the following factors[3]:

- The full flight route occurs within the date range (in this case, between June 1 and June 20).
- The flight route must arrive and leave at least once for NYC, Paris, London, and Tokyo.
- The flight route is valid and logically consistent. For instance, if the first flight takes him from NYC to Tokyo, the second flight in the route should not depart from a place other than Tokyo.

There are a vast pool of valid flight routes. We might assume that out of all these, one would want the *cheapest* route. Of course, this glosses over how the user may want to spend some number of days at each city, but this is the most general case.

Unfortunately, finding the cheapest route possible is challenging using current agencies, because they require knowing the arrival and departure dates for each city. The reason for this is simple: the problem of finding this flight route is a hard problem to solve. In fact, this problem is very similar to the well-known Traveling Salesman Problem (Applegate et al., 2007), which asks if there exists a cycle in a graph that touches each node exactly once (i.e., a Hamiltonian cycle). Each edge has a cost associated with it, and the decision version of the problem, which asks if there exists a valid Hamiltonian cycle with cost bounded by B, is NP-complete (Karp, 1972). In our case, we do not force

---

[3] For brevity, we list a high-level overview of the problem without getting into too many technical details; Section 6 describes some of our assumptions and limitations.

Flight

○ Roundtrip   ○ One way   ⦿ Multiple Destinations

Flight 1: Leaving from: Departing: Time:
[   ]   [mm/dd/yy]   [Any   ⬍]
Going to:
[   ]

Flight 2: Leaving from: Departing: Time:
[   ]   [mm/dd/yy]   [Any   ⬍]
Going to:
[   ]

Flight 3: Leaving from: Departing: Time:
[   ]   [mm/dd/yy]   [Any   ⬍]
Going to:
[   ]

Figure 1. Travelocity lets users select multiple destinations.

the user to visit each city exactly once because the cheapest flight through a sequence of cities *may require* visiting a city more than once.

What makes our problem (likely) much harder than the Traveling Salesman Problem is that the cost of an edge between two nodes (i.e., two cities) is not constant, because flight prices frequently change. We coin our problem the "Distributed Retired Traveling Salesman Problem." "Retired" comes from how we assume that whoever is planning this trip is retired, because otherwise, how would he or she have the time and money? "Distributed" comes from how we incorporate concepts from the design and practice of distributed systems to build software that can tackle this problem. We discuss the mathematical and systems portions of our solution in Sections 3 and 4, respectively.

2. Related Work

The Traveling Salesman Problem (TSP) is one of the oldest and most well-known problems in combinatorial optimization. While no exact, polynomial time solution for the decision problem is known, there is a $3/2$-approximation algorithm due to Christofides (Christofides, 1976). That algorithm was long the standard approximation algorithm until a stunning recent result in (Gharan et al., 2011), which has thus spurred a cascade of additional research relating to the TSP (e.g., (Moemke & Svensson, 2011)).

In contrast to the attention devoted to TSP, there appears to be very little research focusing on the special case of our problem, which introduces additional challenges. These can be mathematically-oriented, such as how to manage the variable costs, or systems-oriented, such as how to even obtain the actual flight costs. We have found nothing in the literature that particularly fits our problem.

3. Mathematical Component

The mathematical portion of our work uses a technique known as *binary integer linear programming*, which falls under the broader category of optimization techniques. The goal in optimization is straightforward: given a set of variables and a set of constraints, the goal is to find the best, feasible solution according to some criteria, such as cost (in which case, "best" means "minimal").

3.1. Linear and Integer Programming

One of the most commonly-used optimization techniques is linear programming, introduced by Dantzig in (Dantzig, 1963).

Definition 3.1. *A* linear programming problem *consists of three components:*

1. *a finite collection of linear inequalities or equations in a finite number of unknowns, $x_1, \ldots, x_n$;*
2. *sign constraints $x_i \geq 0$ on some (possibly empty) subset of the unknowns;*
3. *a linear function to be minimized or maximized.*

An assignment to the variables $x_1, \ldots, x_n$ satisfying the first two conditions is a feasible *solution. If it also satisfies the third, then it is an* optimal *solution (Franklin, 2002).*

Linear programming has tremendous applications, and a full list of them would be impossible to create. Some problems well-suited to linear programming include investment management, scheduling problems, and the diet problem. (The last problem concerns the rather interesting question of what is the minimum cost of a nutritionally adequate diet.) The reason for linear programming's great versatility is the ease at which constraints can be added to a model.

More specific cases of linear programming are integer and binary linear programming. (Sometimes we drop the "linear" part for brevity.)

Definition 3.2. *An* integer programming problem *is a linear programming problem with the added restriction that all variables (i.e., unknowns) $x_1, \ldots, x_n$ are integers.*

Definition 3.3. *A* binary integer programming problem *is an integer programming problem with the added restriction that all variables $x_1, \ldots, x_n$ are such that $x_i \in \{0, 1\}$.*

We see that our problem is particularly suited to binary integer programming, because all the possible flights we could take can be viewed as a set of binary random variables, each of which has value 1 if we decide to take that flight, and 0 otherwise. We now formulate this problem in more detail.

3.2. An Integer Programming Formulation

To start formulating the problem, we make it concrete what we mean by cities, days, and costs.

- We have n cities to reach; we index cities by i or j, where $i, j \in \{1, 2, \ldots, n\}$.
- We have t consecutive days when we can travel: $t \in \{1, 2, \ldots, m\}$.
- The minimum cost for traveling from city i to j on day t is c_{ijt}.

Using the above lets us define the following variables:

$$x_{ijt} = \begin{cases} 1 & \text{if we go from cities } i \text{ to } j \text{ on day } t, \\ 0 & \text{otherwise.} \end{cases}$$

Thus, a solution to our problem will be the full assignment to each of the above variables. The set of the binary variables that equal one will (assuming the set has cardinality k) correspond to k flights $f_1, \ldots, f_k$ ordered by date.

For now, we make several simplifying assumptions, and defer a more detailed discussion about the realism of this project in Section 6. Perhaps the most important one is that we assume each flight is assigned to exactly one day. We also force a valid solution to have flights on unique days. This means we do not consider issues related to (1) overnight flights (we will pretend they do not exist), (2) multiple-stop flights on the same day, and (3) possible logistic impossibilities such as flight f_i arriving at 11:59 PM and then the next flight f_{i+1} departing from the same airport two minutes later (but on the next day).

The goal is to solve this minimization problem:

$$\text{Minimize } \sum_{t=1}^{m} \sum_{i=1}^{n} \sum_{j=1}^{n} c_{ijt} x_{ijt}, \tag{1}$$

subject to the following constraints:

$$\sum_{i=1}^{n} \sum_{j=1}^{n} x_{ijt} \leq 1 \text{ for all } t \in \{1, 2, \ldots, m\}, \tag{2}$$

$$\sum_{t=1}^{m} \sum_{i=1}^{n} x_{ijt} \geq 1 \text{ for all } j \in \{1, 2, \ldots, n\}, \tag{3}$$

$$\sum_{t=1}^{m} \sum_{j=1}^{n} x_{ijt} \geq 1 \text{ for all } i \in \{1, 2, \ldots, n\}. \tag{4}$$

The first constraint ensures that we have at most one flight per day, which is generally too restrictive for real-life, but will suffice for now. The second constraint ensures we enter each city at least once. The third constraint ensures we


Figure 2. This represents two sets of disjoint cycles.

leave each city at least once. These last two reinforce the notion that going through the cheapest route may mean visiting several cities more than once.

Sadly, these previous constraints are *not* enough to solve our problem. If we were to use them, we could get a logically inconsistent flight route due to two main reasons.

3.2.1. DISJOINT CYCLES

The first problem is that the constraints do not prevent disjoint cycles. Figure 2 shows a graph with six nodes, which may represent six cities, and a possible edge assignment that is indicative of the route we take to visit all cities. Assuming the flights are all on different days, the route satisfies our constraints, because we enter and exit each city at least once, but this is not a Hamiltonian cycle.

To prevent disjoint cycles, we need to add constraints that correspond to all the ways we can subdivide the cities into two groups so that both contain at least two cities. (A subdivision where one group only has one city is already covered by our earlier constraints since we would have to leave and enter that city at least once, which requires connecting with the other group.) For instance, with the six-variable situation in Figure 2, one constraint would be

$$\sum_{t=1}^{m} (x_{14t} + x_{15t} + x_{16t} + x_{24t} + x_{25t} + \tag{5}$$
$$x_{26t} + x_{34t} + x_{35t} + x_{36t}) \geq 1,$$

which ensures that at least one leg of the tour connects cities 1, 2, and 3 with cities 4, 5, and 6, so this corresponds to the subdivision of $\{[1, 2, 3], [4, 5, 6]\}$. We can characterize the number of constraints we need to add to prevent disjoint cycles.

Lemma 3.4. *In a problem that has $n \geq 4$ constraints, the number of ways to subdivide the cities so that each has two groups is $\binom{n}{2} + \binom{n}{3} + \cdots + \binom{n}{n-2}$.*

The reasoning is straightforward: we have two groups, and we want to pick the number of elements for one group. We can pick some number out of n elements to be in one group, and assign all the rest to the others.

One technical note is that the number of constraints from

Lemma 3.4 can be halved by symmetry. Still, the number of equations needed grows rapidly with respect to n, and the need to implement these is why integer programming is a hard task. Integer programming is, in fact, an NP-hard problem, and the binary case is NP-complete (Karp, 1972). In contrast, faster algorithms such as the Simplex method[4] are used to solve linear programming problems. Note that simply taking a solution to the linear programming problem and then rounding it to form a "solution" to the integer problem will not work, as we show in Appendix A.

3.2.2. CONSECUTIVE FLIGHTS IN DIFFERENT CITIES

The second problem with the constraints in Section 3.2, which the constraints in 3.2.1 do *not* resolve, is that we can get flight orderings that are logically inconsistent in the sense that consecutive flights may not agree in their choice of cities. It makes no sense, for instance, to have flights f_i and f_{i+1} where f_i takes the participant from NYC to Tokyo, and f_{i+1} takes him or her from Paris to London.

Unfortunately, fixing the flight ordering by adding in constraints is challenging and requires a large number of equations. Our implementation only adds in one layer of constraints for here, and defers a complete flight logic check to the code that actually generates the tree. Specifically, for all pairs of days t and cities c such that $t \in \{1, 2, \ldots, m-1\}$ and $c \in \{1, 2, \ldots, n\}$ we add in the following constraints:

$$\underbrace{\left(\sum_{i \neq c} x_{ict}\right)}_{\text{Term 1}} + \underbrace{\left(\sum_{j \neq c} \sum_{k \neq j} x_{jkt^+}\right)}_{\text{Term 2}} \leq 1, \qquad (6)$$

where t^+ is shorthand for $t + 1$ (which explains why we don't set t to be m). Term 1 represents entering city c, and Term 2 represents leaving any city other than c. Notice that by our previous constraints, both Term 1 and Term 2 are already bounded by one, so if both are one, then this indicates that two flights on back-to-back days are not logically consistent. We therefore have the following lemma.

Lemma 3.5. *Assuming that we have n cities and n or $n + 1$ days as input, solving the binary integer programming problem using the previous constraints will return a valid, logically consistent solution.*

The correctness of n days in the flight schedule will force correctness in the case of $n + 1$ days, but beyond that it is possible to get a logically invalid flight scheduling. We resolve that issue by not using constraints at all. We just

---

[4]The Simplex Method performs well in practice but its runtime has historically been difficult to evaluate, as in the worst case it *can* run in exponential time (Klee & Minty, 1972). The simplex algorithm has *smoothed complexity* polynomial in the input size; for details about this, see (Spielman & Teng, 2004).


Figure 3. This is our ideal machine setup.

run the problem and, each time a candidate solution is proposed, check its complete vector for logical consistency. Thus, it is a "last-minute" check that must be passed before a complete assignment to the variables is accepted.

3.3. Using Balas' Additive Algorithm

To actually solve the binary integer programming problem, we use Balas' Additive Algorithm (Balas, 1965). This is a special case of the larger class of branch-and-bound algorithms, which are used to solve integer programming problems. Branch-and-bound algorithms consist of a systematic enumeration of all candidate solutions, where large subsets of fruitless candidates are discarded by using upper and lower estimated bounds of the quantity being optimized (Clausen, 1997).

Balas' algorithm makes use of the special properties of binary integer programming problems. Assuming that the costs are all non-negative, and that variables are indexed according to cost, we want to (1) set all variables to zero in order to minimize total cost, and (2) if we cannot set all variables to zero due to constraints, then we wish to set the variable with smallest index to one. What results is a rooted tree where each path of n nodes from the root indicates an assignment to the first n variables. One can use a standard depth-first search to implement the algorithm. For brevity we elide additional technical details of the algorithm and refer the reader to Appendix B.

4. Systems Component

We now step back from the mathematics and discuss the implementation. Our code that solves the Distributed Retired Traveling Salesman Problem is made up of several components that form a multi-tiered client-server model (for an introduction to the client-server model, see (Tanenbaum & Steen, 2006)). Figure 3 gives a rough sketch of our ideal machine setup. We have a front-end server, a master server, and a group of slave ("crawler") servers.

```
*************************** DANIEL AND LUCKY'S FLIGHT SCHEDULER ***************************

Instructions: find a series of flights by listing the start date, end date, and cities.

Specific formatting requirements:
        (1) Start and end dates must be the first and second arguments, respectively
        (2) Separate all arguments by at least one whitespace
        (3) Use MM/DD/YYYY format for dates
        (4) Spell city names (or 3-letter abbreviation) correctly
        (5) Use at least three cities

Example request: "06/11/2014 06/16/2014 CHI BOS SEA"
This searches flights from June 11 to 16 (in 2014) that touch Chicago, Boston, and Seattle.

Optional extensions to add at the end of the request:
        (1) Put a number to indicate the minimum days between flights

To exit, type in "q".

*****************************************************************************************

>>> 06/01/2014 06/03/2014 ATL ORD LAX
Input is in the correct format. Now solving ...

['ORD->LAX on 06/01/2014, $189', 'LAX->ATL on 06/02/2014, $206', 'ATL->ORD on 06/03/2014, $190']
```

Figure 4. This shows our front-end server.

4.1. Front-end Server

The front-end server acts as a point of communication between the client[5] and our master server. It is a Python script that starts on the command line and outputs messages to the client to let him or her know how to use our system. The front-end serve receives user input, forwards the request to the master server (step 1 in Figure 3), listens to the master for the result (step 6), and displays it back when ready. Currently, the result is just the single best list of the flights ordered by departure date. If there is more than one optimal solution, it only returns one of them.

Figure 4 shows our current front-end server. Here, the hypothetical user made a request to find the cheapest flight route that touches Atlanta, Chicago, and Los Angeles in the three day period of June 1 through June 3. Our code ran and returned the (valid) flight route of Chicago to Los Angeles to Atlanta that costs $585.

Notice that the front-end server has an optional argument where the user can indicate the minimum number of days between two consecutive flights. (The example from Figure 4 just happened not to use it.) It is straightforward to add more of these optional arguments. If time permits, we also plan to transform the command-line interface to a website, which will expand our audience of potential clients.

4.2. Master Server

The master server is the workhorse of our system. It takes in the user request sent in from the front-end server from Section 4.1. When a request comes in, the master server runs a variety of checks to ensure the user's input is acceptable. For instance, given the optional argument where the user can specify the minimum number of days between flights, the server checks that the full date range is large

---

[5]Somewhat impolitely, we refer to the set of clients, minus Daniel Seita, as the "ignorant masses."

enough to allow that case to happen. We do not catch all possible cases of pathological inputs, in part because we trust our users not to wreck the system. (It is interesting to debate about the impact of trust in systems; for now, we gloss over these and refer the reader to sources such as (Blaze et al., 1999).)

After parsing the input, the master server has two important tasks. The first is that it needs to obtain the cheapest flight that goes from cities i to j on day t for all $n(n-1)m$ possible combinations. It accomplishes this by calling any idle slave servers to look up a single price, which spawns a new thread. Thus, with k total flights and ℓ slaves, we expect each to contribute k/ℓ flight prices to the master. This is shown in Figure 3 as 2 and 4. (That we only have two calls to each slave is just for simplicity, since there can be an arbitrary amount of them.)

Once the master server has all the prices for the flights, it combines this information with the rest of the user request into inputs suitable for binary integer programming. Computing the cost vector is easy; the challenging part is making sure that the constraints are well-defined. Once the master sever has determined the constraint matrix, it calls the binary integer programming solver. As discussed in Section 3.3, this runs Balas' Additive Algorithm to create the final vector of variable assignments. There, if $x_{ijt} = 1$, then the cheapest valid route involves that corresponding flight (and 0 otherwise). It is theoretically possible to distribute Balas' algorithm, but for ease of implementation, we currently run it on one machine, so it is essentially part of the master server. Upon completion, the master server feeds the result back to the client server.

4.3. Slave Servers

The slave servers come from our own machines. Originally, we wanted to use PlanetLab (Peterson et al., 2006), which is a global platform for deploying and evaluating network services, but it was too complicated to work with. Currently, our system can have an unbounded amount of slave servers, each of which must be started before the user sends a request. The slave server's purpose is to, upon request from the master server, use a web crawler to obtain prices of real airlines and relay it back (steps 3 and 5 in Figure 3, though again, the number is arbitrary).

To look up the prices, we implemented a web crawler that extracts prices from the Matrix Airfare Search[6] database, which is powered by ITA Software and is the best flight database currently available. This induces a natural delay in searching for flights because the website must be loaded each time for a search, so for small inputs, the flight lookup process is typically much slower than Balas' algorithm.

---

[6]http://matrix.itasoftware.com/

Table 1. Statistics of our single-crawler system.

INPUT	$|V|$	F TIME	NODES	B TIME
3C 3D	18	02:31	174	0:00
3C 4D	24	04:01	717	0:00
3C 5D	30	04:23	2845	0:00
3C 6D	36	05:14	10336	0:00
3C 7D	42	06:03	35902	0:00
3C 8D	48	06:47	127570	0:01
3C 9D	54	07:36	460872	0:08
3C 10D	60	08:45	1472853	0:32
3C 11D	66	09:22	4609309	2:23
3C 5D 1M	30	04:23	2855	0:00
3C 6D 1M	36	05:29	10136	0:00
3C 7D 1M	42	06:07	38613	0:00
3C 8D 1M	48	06:41	125988	0:01
3C 9D 1M	54	07:38	442564	0:08
3C 7D 2M	42	06:23	38217	0:00
3C 8D 2M	48	06:53	125267	0:01
3C 9D 2M	54	07:53	442122	0:08
3C 9D 3M	54	07:48	442412	0:09
4C 7D	84	12:21	1548120	0:45
5C 5D	100	14:14	261089	0:10

Table 2. Our system with multiple crawlers.

INPUT	THREADS	F TIME
3C 12D	1	TODO
3C 12D	2	05:01
3C 12D	3	03:20
3C 12D	4	02:38
3C 12D	5	02:09
3C 12D	6	01:51

5. Experiments and Results

We present experimental results and analyze our system. Due to space constraints, we only offer limited descriptions. Appendix C contains more complete results and includes additional statistics not discussed here.

5.1. Single-Node Experiments

We ran our system using a variety of inputs on a single node (i.e., crawler). Table 1 shows our results on twenty inputs. The inputs have Cs, Ds, and Ms, denoting the number of cities, days, and minimum days between flights (if any). The cities were Atlanta, Chicago, and Los Angeles. The main concern is the runtime of our whole system. The "F Time" and "B Time" columns represent the time (minutes:seconds) to look up flight prices and run Balas' algorithm, respectively. "V" represents the set of variables, and the "Nodes" column is how many nodes Balas' algorithm expands. In order to be used in practice, the number of nodes _must be substantially smaller_ than the full list of $2^{|V|}$ possible vector assignments.

From the results, we see that in terms of $|V|$, flight time increases linearly but Balas' time increases exponentially (as expected). We tried running the following inputs: 3C 12D, 4C 8D, and 5C 7D, but for each one, Balas' algorithm would continue for hours without making progress. Once we start increasing the number of days and cities beyond the rough values shown in Table 1, Balas' algorithm will be the limiting factor of our system.

5.2. Multi-Node Experiments

The goal here is to understand the impact of adding in multiple crawler servers on the flight time. Ideally, the runtime of flight lookup should decrease with respect to the number of threads, but each thread added results in a relatively smaller decrease in time (i.e., diminishing returns). We can also gain insight about the threshold when Balas' algorithm starts becoming more expensive than flight checking by comparing the runtime of the same request on different number of crawler servers. And if Balas' algorithm is not the limiting factor of the algorithm, we can always add more threads.

Table 2 shows our results (times are in mintes:seconds format), which confirm our hypothesis. Thus, for even moderately large inputs, flight checking can be done relatively quickly, and extra effort should be devoted towards distributing Balas' algorithm, or finding another way to solve integer programming.

6. Assumptions and Limitations

This section addresses a number of issues briefly touched upon in Section 3.2 and other areas of this paper. We avoid discussing anything that can be easily added as an extension (e.g., forcing us to start at one particular city, or making the flight route to make a stop at a city three times).

6.1. Practical Usage

An obvious question to consider is whether our system is practical for real-life use. The experiments in Section 5 show that it is doable for a small enough variable vector. The values in Table 1 are a good guideline as to what is acceptable, but keep in mind that those were done on just one machine. In a real "business-like" situation, this kind of algorithm would be distributed over hundreds of machines. Then again, it is unlikely that a real vacation would span, for instance, twenty cities over sixty days. As a general guideline, letting n be the number of cities and m be the number of days, we recommend bounding $n(n-1)m$, or the length of the variable vector, by 100.

To speed up the algorithm, we can easily distribute the flight lookup process (e.g., using multiple threads). Unfortunately, as that process runs quicker, Balas' algorithm will become the major time bottleneck. With a large variable vector, the depth-first search process takes prohibitively long since the number of paths to explore is exponential. One workaround would be to add in more constraints, so that we can prune more often, but unless the constraints are overly restrictive — which defeats our objective of having the user supply minimal information — the number of nodes to search will still grow quickly based on $n(n-1)m$. Balas' algorithm can be distributed by assuming that the vector starts with a certain value and then assigning the algorithm to determine the rest of them. For instance, with four machines, we can have four computers each solve Balas' algorithm under the assumption that their vectors start with $[0,0]$, $[0,1]$, $[1,0]$, and $[1,1]$, respectively. Due to the exponential time complexity of Balas' algorithm, this will not be a scalable solution, but may work if we limit the user to requesting "reasonable" flight routes.

6.2. One Flight a Day

As stated earlier, one of our key assumptions is that every flight lasts one day and each day can only have one flight. In real life, however, it is common for people to take two, or even three, flights in a day, especially if their departure or destination cities are not one of the major airline hubs (e.g., Chicago O'Hare or Hartsfield-Jackson Atlanta).

This limitation is, in fact, often not a problem. The key is that when looking up flights between two cities, Matrix ITA defaults to searching not only for direct flights, but also for flight routes that make one or two stops. So when making a request to go from Albany to Los Angeles, our crawler will pick the cheapest flight that appears on the list, which almost always makes one stop at airports such as Chicago or Charlotte. So even though our code might output `[ALB->LAX, 07/07/2014, $400]`, there will be another flight "hidden" in our route.

A more problematic situation with our "one flight" a day limitation is that it may be cheaper to take two separate flights that touch three major cities. For example, if we wanted to travel to Chicago, Atlanta, and Denver in some order over the span of six days, the cheapest route may involve going from Chicago to Atlanta on the third day, and then going from Atlanta to Denver, *also* on the third day.

Another limitation is that we only view time in terms of absolute days. If we look up a flight on day t, an overnight flights that starts on day t may be selected, and we the code may also assign a flight to start on day $t + 1$, raising some logistical issues. Even if we ignore overnight flights, we could have a flight landing at the very end of day t and the next flight departing at the beginning of day $t + 1$ with

insufficient layover time. We do not have ways of facing this limitation in our system, but one possible idea is to treat time in one-hour units, rather than 24-hour units, giving us a finer grain of control. It is unclear how much this would complicate the rest of our constraints.

6.3. One Ticket, Multiple Flights

One issue related to optimality is that even if we search for all possible direct flight combinations and find the cheapest collective route, it may still not be the best possible route. Consider a hypothetical situation where a user wants to travel to Boston, Chicago, London, and Beijing. Analyzing all pairwise combinations might result in the following route: Chicago to Boston, Boston to London, London to Beijing, then Beijing to Chicago. Assume for simplicity that each flight costs $500. But there might exist a flight from Chicago to London that, according to the Matrix Airfare Search, costs $750, *but makes a stop at Boston*, despite how the cheapest direct flight from Chicago to Boston costs $500 (and similarly for Boston and London). In other words, a better route would be Chicago to London (making a stop at Boston), then London to Beijing, then Beijing to Chicago, which would cost $1750, compared to the $2000 route that our solution would provide. To resolve this, we are looking to expand our crawler so that it can detect the stopgap cities that are part of a single flight.

6.4. Web Crawler Issues

In general, our crawler is safe enough that we are confident in running it and only periodically checking back on its results. One problem, though, is that it sometimes gets stuck on a Matrix Airfare Search page after searching the price for a flight. Typically, this happens when we have it running "in the background" (in OS X, this problem can happen when one switches between windows — we have not tested on other operating systems); bringing it back up to the foreground successfully makes it continue. There are also some problems with foreign currency units. For now, we assume that the person using this service will originate from the United States, which forces money to be in dollars. Finally, in rare cases there may be no flight found between two cities, though each time we ran into this problem, it was due to poorly formatted input.

7. Conclusion

We view this as a beginning, rather than an end, of the Distributed Retired Traveling Salesman Problem. There are several ways of proceeding from our work.

- In terms of the mathematical component, there may be other valid ways of solving binary integer programming problems, and it would be of interest to compare

the feasibility of different algorithms.

- We could address the limitations we discussed in Section 6. The biggest one might be to extend the "time" constraint so that we can always ensure that our flight route has sufficient layover time at airports.
- We could improve the aesthetics of our implementation. Most prominently, our front-end server is in the form of a command-line interface, but many people do not know how to use those. Consequently, we could reach out to and would be much more comfortable with a streamlined website.

The most important overall goal moving forward should be to figure out how to improve the system's runtime. While the algorithm will return "correct" flight routes, the runtime can be prohibitive for even a moderate amount of cities and days. This can be mitigated by distributing price lookups and *especially* distributing Balas' algorithm, but the latter will continue to be a major time bottleneck. Striking the right balance between piling on constraints/limitations (decreases runtime, increases restrictiveness) and brute-force searching (increases runtime, decreases restrictiveness) seems to be the way to go to provide a system that allows the user to plan out a cheap trip without checking all possible flight routes.

Acknowledgments

References

Applegate, David L., Bixby, Robert E., Chvatal, Vasek, and Cook, William J. *The Traveling Salesman Problem: A Computational Study (Princeton Series in Applied Mathematics)*. Princeton University Press, Princeton, NJ, USA, 2007. ISBN 0691129932, 9780691129938.

Balas, Egon. An additive algorithm for solving linear programs with zero-one variables. *Operations Research*, 13 (4):517–546, 1965. doi: 10.1287/opre.13.4.517.

Blaze, Matt, Feigenbaum, Joan, Ioannidis, John, and Keromytis, Angelos D. Secure internet programming. chapter The Role of Trust Management in Distributed Systems Security, pp. 185–210. Springer-Verlag, London, UK, UK, 1999. ISBN 3-540-66130-1.

Bradley, S.P., Hax, A.C., and Magnanti, T.L. *Applied mathematical programming*. Addison-Wesley Pub. Co., 1977. ISBN 9780201004649.

Chinnek, John. Practical optimization: A gentle introduction. Online Textbook, 2014.

Christofides, Nicos. Worst-case analysis of a new heuristic for the travelling salesman problem. Technical Report 388, Graduate School of Industrial Administration, Carnegie Mellon University, 1976.

Clausen, Jens. Branch and bound algorithms-principles and examples. *Parallel Computing in Optimization*, pp. 239–267, 1997.

Dantzig, George B. *Linear programming and extensions*. Rand Corporation Research Study. Princeton Univ. Press, Princeton, NJ, 1963.

Franklin, Joel N. *Methods of mathematical economics : linear and nonlinear programming, fixed-point theorems*. Classics in applied mathematics. SIAM, Philadelphia, 2002. ISBN 0-89871-509-1.

Gharan, Shayan Oveis, Saberi, Amin, and Singh, Mohit. A randomized rounding approach to the traveling salesman problem. In Ostrovsky, Rafail (ed.), *FOCS*, pp. 550–559. IEEE, 2011. ISBN 978-1-4577-1843-4.

Karp, R. Reducibility among combinatorial problems. In Miller, R. and Thatcher, J. (eds.), *Complexity of Computer Computations*, pp. 85–103. Plenum Press, 1972.

Klee, V. and Minty, G. J. How Good is the Simplex Algorithm? In Shisha, O. (ed.), *Inequalities III*, pp. 159–175. Academic Press Inc., New York, 1972.

Miller, Steven. An introduction to advanced linear algebra. Operations Research Class Notes, 2012.

Moemke, Tobias and Svensson, Ola. Approximating graphic tsp by matchings. In *Proceedings of the 2011 IEEE 52Nd Annual Symposium on Foundations of Computer Science*, FOCS '11, pp. 560–569, Washington, DC, USA, 2011. IEEE Computer Society. ISBN 978-0-7695-4571-4. doi: 10.1109/FOCS.2011.56.

Peterson, Larry L., Bavier, Andy C., Fiuczynski, Marc E., and Muir, Steve. Experiences building planetlab. In Bershad, Brian N. and Mogul, Jeffrey C. (eds.), *OSDI*, pp. 351–366. USENIX Association, 2006. ISBN 1-931971-47-1.

Spielman, Daniel A. and Teng, Shang-Hua. Smoothed analysis of algorithms: Why the simplex algorithm usually takes polynomial time. *J. ACM*, 51(3):385–463, May 2004. ISSN 0004-5411. doi: 10.1145/990308.990310.

Tanenbaum, Andrew S. and Steen, Maarten van. *Distributed Systems: Principles and Paradigms (2Nd Edition)*. Prentice-Hall, Inc., Upper Saddle River, NJ, USA, 2006. ISBN 0132392275.

Supplementary Material

Outline of Supplementary Material:

- why solutions to linear and integer programming can differ enormously
- a detailed discussion of Balas' Additive Algorithm and our implementation
- additional experiments and examples that could not fit in the text

We deemed the first eight pages of this document sufficient for presenting our results. The supplementary material is for anyone who really wants to *get* what we did, as we glossed over a few details for brevity.

A. Integer Programming versus Linear Programming

The Simplex method can usually solve linear programming problems quickly. Unfortunately, there is no efficient analogue of the Simplex method (or similar algorithms) to the case when all variables must be integers. To make matters worse, a problem can have optimal integral and linear solutions, but they may be extraordinarily different, so it is not sufficient to simply round, truncate, or investigate the "relative area" surrounding a linear programming solution. The following problem provides a concrete example and is based on the presentation of (Miller, 2012).

Problem A.1. *[The Knapsack Problem] Imagine we have a knapsack that can hold at most 100 kilograms. There are three items we can pack, each of which is worth some monetary amount per unit, and the goal is to carry as much value in the knapsack as possible. The first item weighs 51 kilograms and is worth \$150 per unit. The second item weights 50 kilograms and is worth \$100 per unit. Finally, the third item also weighs 50 kilograms, but is worth \$99 per unit. If we let x_1, x_2 and x_3 represent the amount of the first, second, and third items, respectively, and we assume that we can't take on any negative quantity, the constraint for our problem is*

$$51x_1 + 50x_2 + 50x_3 \leq 100, \tag{7}$$

and we want to maximize

$$150x_1 + 100x_2 + 99x_3. \tag{8}$$

Before presenting the solution for both the integral and linear case, think about the relative "bang of the buck" for the items. The second and third items both weigh the same, but the second is worth a dollar more, so clearly, we prefer it to the third. The first weighs just a fraction more than the second (51 versus 50 kilograms), but is worth 50 percent more in dollars (\$150 to \$100), so clearly it's the best value item. If we allow all x_j to be real numbers, the optimal answer is $x_1 = 100/51$ and $x_2 = x_3 = 0$; the value of the knapsack is about \$294.12. This makes sense, given our previous discussion. If we require all x_j to be integers, however, the optimal solution is $x_2 = 2$ and $x_1 = x_3 = 0$. The value of the knapsack is \$200, significantly lower than the linear case since the 51 kilogram weight of the first item just makes it ineligible for us to have two of it.

Incidentally, Problem A.1 demonstrates how, assuming the objective is maximization, the linear programming solution is an upper bound to the integer programming solution (and lower bound if the goal is to minimize something).

Another enlightening example of the difference between integer and linear programming is in Chapter 9 of (Bradley et al., 1977), which also demonstrates how the optimal linear and integral solutions can be arbitrarily far away from each other, in terms of distance in the plane or total cost.

B. Balas' Additive Algorithm

We briefly introduced Balas' Additive Algorithm in Section 3.3 for solving binary integer programming problems, and in this section we present a deeper treatment. Our presentation is heavily influenced by (Chinnek, 2014).

B.1. Canonical Form

Balas' algorithm first requires the input to be converted to canonical form, if it is not done so already. Letting the variables and their corresponding costs be $x_1, \ldots, x_n$ and $c_1, \ldots, c_n$, respectively, canonical form is when the following are satisfied:

- The objective function is to minimize $\sum_{j=1}^{n} c_j x_j$.
- All objective function coefficients (i.e., variable costs) are non-negative.
- The m constraints are all inequalities of the form $\sum_{j=1}^{n} a_{ij} x_j \geq b_i$ for $i \in \{1, 2, \ldots, m\}$.
- All variables are ordered according to their costs, so for $x_1, x_2, \ldots, x_n$, we have $0 \leq c_1 \leq c_2 \leq \cdots \leq c_n$.

Incidentally, our code to solve Balas' algorithm *does* include checking and conversion capability, which is necessary for some constraints, such as the one requiring at most one flight per day. We also order variables by cost, so any final vector of, say, $[1, 0, 0, 1, 0, 0, 1]$ means that we pick the three flights that correspond to the cheapest, fourth cheapest, and seventh cheapest (i.e., the most expensive one).

While it may seem like requiring the input to be in canonical form is a heavy restriction, in fact it is relatively simple to do the conversion. Ordering the constraints should not be challenging in an implementation, given enough debugging to ensure that indices and other aspects line up correctly. With a less than or equal bound in one of the constraints, we multiply through by -1. With negative cost coefficients, a change of variables is required, with transforming x_i into $1 - x_i'$. Finally, to handle equality constraints, we can transform the equality into two inequalities, similar to how we can convert $x = 2$ to the equivalent condition of $x \leq 2$ and $x \geq 2$ (and then we would subsequently change the first one to be $-x \geq 2$).

As we mention in the text, the idea of Balas' algorithm is that we want to assign as many variables to zero as possible since we know that coefficients are non-negative. In almost all cases, if we tried to do this at the start, we would violate one of the constraints. So the next step is to try and set the *cheapest* variables to be one, because if we can satisfy the constraints, we prefer to do it with the cheapest variables possible.

Our implementation follows a depth-first search algorithm to enumerate all possible solutions. In other words, we form a rooted tree where a path corresponds to some assignment of variables, and each node has two successors in the tree, corresponding to a selection of 0 or a selection of 1 for the current variable, where the "current" variable is determined by the level of the tree (i.e., the root node corresponds to an empty path $[\ldots]$[7], its two successors correspond to the paths of $[1, \ldots]$ and $[0, \ldots]$, and so on).

But a word of caution — if there are four cities and twenty days to travel, that means we have $4 \cdot 3 \cdot 10 = 120$ total variables, which means listing all solutions would mean we have 2^{120} leaf nodes. This is clearly impractical, and the hope is that Balas' algorithm can avoid searching the full space of solutions by using look-ahead and pruning techniques, which we now describe.

B.2. Cost and Feasibility Look-Ahead

Suppose we are at a given node of the DFS tree with path $[x_1, x_2, \ldots, x_j, \ldots]$, so the first j variables are assigned and the remaining $n - j$ variables are unknown. It is possible to "look ahead" in the tree to see if best-case (i.e., lowest-cost) scenarios are feasible. We can consider two cases for the value of x_j:

- Suppose $x_j = 1$. Then we need to check if the current path so far can lead to a best-cost solution. If we check the constraints, and find that the assignment $\mathcal{X} = [x_1, x_2, \ldots, x_j = 1, 0, 0, \ldots, 0]$ works (i.e., at least the last $n - j$ variables are zero), then we compute the cost of $\mathcal{X}$ and compare it to the current best cost known. If $\mathcal{X}$ is better than the current best cost, then we save its path and cost. Furthermore, *there is no need to continue in this DFS branch*, because any further path must start with the first j assignments in $\mathcal{X}$, and any extra one is going to force a subsequent solution $\mathcal{X}'$ to be costlier than $\mathcal{X}$. This is the power of knowing that the costs are ordered, and that each variable is binary. Without either of those assumptions, Balas' algorithm would not work!

---

[7]We use the ellipses as shorthand for "the rest of the variables are unknown."

- Now suppose $x_j = 0$. Here, it is important to assume that the assignment $\mathcal{X} = [x_1, x_2, \ldots, x_j = 0, 0, 0, \ldots, 0]$ is *infeasible*. The reasoning is that if it were feasible, then we would not be in this node at all because then the previous node would have already been feasible. Since we are assuming best-case costs each time, this means that the DFS would not have continued after the previous node was expanded. Thus, we must assume $\mathcal{X}$ is infeasible. Then we set the *next* variable $x_{j+1} = 1$, and then after that, check to see if the solution $\mathcal{X}' = [x_1, x_2, \ldots, x_j = 0, x_{j+1} = 1, 0, 0, \ldots, 0]$ is feasible. If it is, there is no need to expand this DFS path any further.

As we will soon show, it is also important to be able to prune away impossible paths so that we do not need to do this "cost look-ahead" every time.

B.3. Pruning

We might have feasible solutions, but it is also important to be able to stop checking a DFS path if we can tell that its best-case flight route price is always going to be more expensive than the current best cost. Given a current path assignment $\mathcal{X} = [x_1, \ldots, x_j, \ldots]$, the way we do this is by analyzing each constraint one by one. The constraints are all of the form $\sum_{j=1}^{n} a_{ij} x_j \geq b_i$ for $i \in \{1, 2, \ldots, m\}$. So the key is to determine the *largest possible* value of the left hand side. If that value is still less than b_i, we can avoid expanding that path any further.

To make the left hand side as large as possible, set each *unknown* variable x_i to be zero if its corresponding coefficient is negative, and one if otherwise. For example, if we had the assignment $[x_1 = 1, x_2 = 0, x_3, x_4]$ (where the last two are unknown) and a constraint $x_1 - 5x_2 + 2x_3 - 3x_4 \geq 4$, then knowing the first two values mean this constraint is immediately $1 + 2x_3 - 3x_4$. Set $x_3 = 1$ and $x_4 = 0$ and we get $1 + 2 \not\geq 4$, so this can be pruned.

B.4. Our Implementation

Our implementation of the algorithm runs an iterative version of depth-first search, where each node contains information about its current path from the root to itself. We debated between depth-first search versus best-first search, but decided to do the former since best-first search would give priority early to paths that start with almost all zeros (and we want paths that start with a few ones and end with lots of zeros).

An interesting detail about our code is that we actually view each node as having three successors in the tree. Given a current path $\mathcal{X} = [x_1, \ldots, x_j]$, we consider the three paths that result from appending $[1], [0, 1]$, or $[0, 0]$ to $\mathcal{X}$. We only run the cost look-ahead procedure (described in Appendix B.2) on the first two, because adding in a bunch of zeros to an infeasible assignment $\mathcal{X}$ cannot make it feasible. Remember, the whole reason why we are still expanding $\mathcal{X}$ is because best-case extensions of it (i.e., adding all zeroes, possibly after adding in a one to start) is not feasible! Our original implementation simply had two children, which were formed from appending $[0]$ and $[1]$. This was inefficient as it caused us to repeat certain trials. After making each node expand three children, we noticed that the depth-first search was expanding a little over *half* as many nodes as before! Any small improvement we can make to this algorithm is imperative, given its general exponential-time complexity with respect to the number of cities and days.

C. Experiments on a Single Machine

In this section, we describe an in-depth demonstration of our system in action on one machine. The goal is to see how varying the parameters causes the solution to change (if it does), and to see how factors such as pruning and look-aheads affect the speed of our depth-first search. Be aware, though, that there is always a risk of experimenting with real flight data, because prices can change at any time[8]. To mitigate this risk, we run these experiments back-to-back. We also decided to use the busiest American airports [9] in our input to avoid any potential complications with missing flights on certain days. (Future work should improve our code's failure capabilities.)

Specifically, we run our flight scheduler code with the following requests. First, we have the ones from Sections C.1 and C.2:

- `06/01/2014 06/03/2014 ATL ORD LAX`
- `06/01/2014 06/04/2014 ATL ORD LAX`
- `06/01/2014 06/05/2014 ATL ORD LAX`
- `06/01/2014 06/06/2014 ATL ORD LAX`
- `06/01/2014 06/07/2014 ATL ORD LAX`
- `06/01/2014 06/08/2014 ATL ORD LAX`
- `06/01/2014 06/09/2014 ATL ORD LAX`
- `06/01/2014 06/10/2014 ATL ORD LAX`
- `06/01/2014 06/11/2014 ATL ORD LAX`

Then we have the ones from Section C.3:

- `06/01/2014 06/05/2014 ATL ORD LAX 1`
- `06/01/2014 06/06/2014 ATL ORD LAX 1`
- `06/01/2014 06/07/2014 ATL ORD LAX 1`
- `06/01/2014 06/08/2014 ATL ORD LAX 1`
- `06/01/2014 06/09/2014 ATL ORD LAX 1`
- `06/01/2014 06/07/2014 ATL ORD LAX 2`
- `06/01/2014 06/08/2014 ATL ORD LAX 2`
- `06/01/2014 06/09/2014 ATL ORD LAX 2`
- `06/01/2014 06/09/2014 ATL ORD LAX 3`

Finally, the two in Section C.4 are[10]:

- `06/01/2014 06/07/2014 ATL ORD LAX MIA`
- `06/01/2014 06/05/2014 ATL ORD LAX MIA DFW`

Notice that a four city, eight day request ran Balas' algorithm for *several hours*, before we decided to quit it because it was slowing down, most likely due to eating up too much memory. It had $4 \cdot 3 \cdot 8 = 96$ variables, and we suspect that 100 is around the time when things start becoming infeasible (at least, on one machine with limited memory). Five cities and six days (with 100 variables) *could* run on one machine, though five cities and seven days probably would not. With three cities, we suspect the feasible limit is around 14 days; after that, it's too much. For six cities, we think six days is the best possible. Seven or more is too much.

We record the following quantities:

- The time it takes to search for the flight prices[11]

---

[8]In fact, a week before finishing this paper, we ran an experiment in which the cost of one flight decreased by \$11 in between our trials, which was why we got such surprising results that day.

[9]http://en.wikipedia.org/wiki/List_of_the_busiest_airports_in_the_United_States

[10]The number of experiments is already quite high, so we basically used the last section to explore the limits of our solver.

[11]As we mention in Section 6, sometimes our web crawler gets stick on Matrix Airfare Search. We resolve this by simply leaving the Matrix Airfare Search webpage open at all times (i.e., we do not minimize it or switch views to a different window).

- The time it takes to complete Balas' Additive Algorithm
- The number of nodes expanded in the depth-first search (DFS) for Balas' algorithm
- The number of pruning checks Balas' algorithm makes, as well as the number of successes
- The number of times Balas' algorithm "looks ahead" by filling out a path and checking for feasibility and cost.
- The number of flight logic errors due to either violating flight ordering, or violating the minimum number of days between flights requirement.

These experiments as a whole therefore test the effect of adding more days, adding more cities, and adding in the extra constraint of the minimum number of days between cities. The goal is to gain insights about the feasibility of using our flight scheduler. Notice that the time it takes for the entire system to complete a request is almost entirely dependent on the flight lookup and Balas' algorithm.

C.1. The First Flight Request

In this section (and some subsequent ones) we print out part of the Master Server's output[12]. We describe output for this first flight request in detail for explanatory purposes, and then we are more concise after this. Here, the input is:

```
06/01/2014 06/03/2014 ATL ORD LAX
```

The output is:

```
Starting the problem! Here's our input:
All dates: [06/01/2014, 06/02/2014, 06/03/2014]
All cities: [ATL, ORD, LAX]

Now checking flight prices...
ATL->ORD on 06/01/2014, $315
ATL->LAX on 06/01/2014, $261
ORD->ATL on 06/01/2014, $315
ORD->LAX on 06/01/2014, $253
LAX->ATL on 06/01/2014, $278
LAX->ORD on 06/01/2014, $204
ATL->ORD on 06/02/2014, $210
ATL->LAX on 06/02/2014, $201
ORD->ATL on 06/02/2014, $210
ORD->LAX on 06/02/2014, $253
LAX->ATL on 06/02/2014, $201
LAX->ORD on 06/02/2014, $174
ATL->ORD on 06/03/2014, $190
ATL->LAX on 06/03/2014, $172
ORD->ATL on 06/03/2014, $190
ORD->LAX on 06/03/2014, $154
LAX->ATL on 06/03/2014, $171
LAX->ORD on 06/03/2014, $154

Now converting to an integer programming problem...
Here is our power set of cities with at least 2 and 2 in the subsets: []
Here are the problem inputs:

The costs: [315, 261, 315, 253, 278, 204, 210, 201, 210, 253, 201,
174, 190, 172, 190, 154, 171, 154], and the 15 constraints:

[1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1]
[0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1]
[0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1]
[1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1]
[0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1]
[1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1]
```

---

[12]This output could, theoretically, be part of the front-end that the user actually sees, but we decided that this is too much of an information overload for them, so we do not bother to do that.

```
[0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1]
[0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1]
[0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1]
[1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1]
[0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1]
[0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1]
[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1]
[0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1]

Now calling Balas' algorithm...

Now doing the depth first search...

Case 2 Update: Node #22, New cost: 642,
New path: [1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0]
Case 1 Update: Node #87, New cost: 586,
New path: [0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0]

Total nodes expanded: 174, total cost/feasibility look-aheads: 348
Total flight logic errors caught: 0, total min-days errors caught: 0
Total pruning checks: 513, total successful: 340
Best solution: [0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0], with cost 586

Original problem from client was 06/01/2014 06/04/2014 ATL ORD LAX
Flights ordered by cost:
[ATL->LAX on 06/03/2014, $172, LAX->ORD on 06/01/2014, $204, ORD->ATL on 06/02/2014, $210]
Flights ordered by date:
[LAX->ORD on 06/01/2014, $204, ORD->ATL on 06/02/2014, $210, ATL->LAX on 06/03/2014, $172]
Flight check time in hours:minutes:seconds -- 0:2:31.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

Let us discuss this output in more detail. The first part that prints out the date range and cities is essentially to confirm that the user's input is well-formed. Next, the Master Server prints the flight tickets for each flight as it gets them by querying the crawler server. (For small enough inputs, looking at these prices can serve as a useful sanity check to make sure that our code actually outputs the correct solution.) Next up, the Master Server prints something about a "power set." This is to solve the disjoint cycle problem mentioned in Section 3.2.1, and lists all the ways to partition the cities into groups. With three cities, this is a pointless exercise because every possible subdivision of three into two groups results in one grouping having either zero or one city, which is already covered by other constraints and thus is not problematic.

Next, the Master Server prints out the costs and the constraints. The costs are straightforward and can be directly verified by the Master Server's printing of the flight information as they get discovered. The constraints are the interesting part. *Most* of the numbers there correspond to coefficients. The exception are the last two numbers in each list. The penultimate number in each list follows a code where a zero indicates *less than* and a one indicates *greater than*[13]. Finally, the last element in each list is the bound.

With the constraints, we can then call Balas' algorithm. Periodically, the algorithm will find a working solution via "look-aheads" and those are the "updates" that are being printed out, along with the corresponding node, the cost, and the (complete) path. Cases 1 and 2 refer to when an update is found via a look-ahead that corresponds to appending a [1] or a [0, 1] to the node's list, respectively. In this example, there was a valid flight route found to have cost $642, but that was later bested by the $586 solution.

Finally, once the algorithm is finished, the server prints some concluding remarks and statistics. Most importantly, it prints out the best flight route, and its total cost. In the case of a simple three city, three day input, the best flight route costs $586 and goes from Los Angeles to Chicago to Atlanta to Los Angeles. The total flight checking time took two minutes and thirty one seconds, and the run time of Balas' algorithm is negligible. The total number of nodes expanded is 174, which is a substantial savings over the 2^{18} complete possible enumerations of the flight vector. As expected, there are no flight logic errors caught (because those require more days than flights) and there are no "minimum day" errors because we did not specify any.

---

[13]Technically, it might have been better to assign a minus-one to mean less than, since then we could just multiply everything by -1 to get a greater than. As it is, we do the multiplication by -1, but have to explicitly set the penultimate element of such a list to be 1.

C.2. Additional Three-City, No Minimum Day Queries

This section describes additional three-city outputs without any required minimum number of days between flights. We will be more concise than in Section C.1 and include just the information we deem essential. We also briefly comment on the output of each experiment. Recall that with the original request in Section C.1, the best flight cost $586 and corresponded to the route:

```
[LAX->ORD on 06/01/2014, $204, ORD->ATL on 06/02/2014, $210, ATL->LAX on 06/03/2014, $172]
```

So what happens when we add in an extra day? The following input:

```
06/01/2014 06/04/2014 ATL ORD LAX
```

has output:

```
Case 2 Update: Node #26, New cost: 814,
New path: [1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0]
Case 1 Update: Node #26, New cost: 777,
New path: [1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0]
Case 1 Update: Node #43, New cost: 761,
New path: [1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
Case 2 Update: Node #47, New cost: 721,
New path: [1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0]
Case 2 Update: Node #51, New cost: 517,
New path: [1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0]
Case 2 Update: Node #187, New cost: 505,
New path: [0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
Total nodes expanded: 717, total cost/feasibility look-aheads: 1434
Total flight logic errors caught: 0, total min-days errors caught: 0
Total pruning checks: 2086, total successful: 1370
Best solution: [0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 505

Original problem from client was 06/01/2014 06/04/2014 ATL ORD LAX
Flights ordered by cost:
[ORD->ATL on 06/04/2014, $150, LAX->ORD on 06/03/2014, $154, ATL->LAX on 06/02/2014, $201]
Flights ordered by date:
[ATL->LAX on 06/02/2014, $201, LAX->ORD on 06/03/2014, $154, ORD->ATL on 06/04/2014, $150]
Flight check time in hours:minutes:seconds -- 0:4:1.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

Interesting! Now the best flight route is in a completely different order on different days (June 2, 3, and 4 compared to June 1, 2, and 3), and costs $505. This is a good sanity check, because unless our flight prices somehow changed between experiments, our new best price should clearly not be more expensive than $586. So the current flight route to beat is:

```
[ATL->LAX on 06/02/2014, $201, LAX->ORD on 06/03/2014, $154, ORD->ATL on 06/04/2014, $150]
```

One more thing: the flight lookup time now takes just a second over four minutes, but Balas' algorithm *still* runs in less than a second. Let's hope that is still the case when we add another day. The following input:

```
06/01/2014 06/05/2014 ATL ORD LAX
```

has output:

```
Case 2 Update: Node #32, New cost: 923,
New path: [1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0]
Case 2 Update: Node #49, New cost: 654,
New path: [1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0]
Case 2 Update: Node #152, New cost: 650,
New path: [1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0]
Case 2 Update: Node #160, New cost: 458,
New path: [1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
Total nodes expanded: 2845, total cost/feasibility look-aheads: 5690
Total flight logic errors caught: 0, total min-days errors caught: 0
Total pruning checks: 8141, total successful: 5297
Best solution:
[1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 458

Original problem from client was 06/01/2014 06/05/2014 ATL ORD LAX
Flights ordered by cost: [LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/03/2014, $172]
Flights ordered by date: [ATL->LAX on 06/03/2014, $172, LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151]
Flight check time in hours:minutes:seconds -- 0:4:23.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

We get *another* improvement! This time, the cost went down to $458, from a route that has flights on the third, fourth, and fifth. (Unsurprisingly, any route that improves on a previous one *must* have one flight on the new day we just added.) Notice again that the runtime of the flight lookup is around four minutes, while Balas' algorithm takes negligible time.

Let us add a day. The following input:

```
06/01/2014 06/06/2014 ATL ORD LAX
```

has output (note: we remove the full paths since they start to get long):

```
Case 2 Update: Node #20, New cost: 1052,
Case 2 Update: Node #57, New cost: 827,
Case 2 Update: Node #137, New cost: 635,
Case 2 Update: Node #425, New cost: 622,
Case 2 Update: Node #447, New cost: 458,
Total nodes expanded: 10336, total cost/feasibility look-aheads: 20672
Total flight logic errors caught: 0, total min-days errors caught: 0
Total pruning checks: 29293, total successful: 18958
Best solution:
[1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
with cost 458

Original problem from client was 06/01/2014 06/06/2014 ATL ORD LAX
Flights ordered by cost: [LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/03/2014, $172]
Flights ordered by date: [ATL->LAX on 06/03/2014, $172, LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151]
Flight check time in hours:minutes:seconds -- 0:5:14.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

Sadly, the streak of improvement was not to be. Despite adding in June 6 as a possible date, the cheapest flight was still the one that spanned June 3 through June 5. On the positive side, flight checking time (five minutes) is still reasonable, especially for one machine. Balas' algorithm continues to be fast, though it is important to note that we're now expanding tens of thousands of nodes (10,336 in this case) and that will start to grow exponentially. Some more interesting observations are that (1) we continue to find most new paths by virtue of Case 2, which is not surprising because the data is generally sparse and most of the time, ones are preceded by zeros[14] and that (2) we still do not have any flight logic errors to catch, and that (3) pruning is successful for about two-thirds of the time, showing that it is very valuable to do so!

With that aside, let us add in *another* day. The following input:

```
06/01/2014 06/07/2014 ATL ORD LAX
```

has output:

```
Case 2 Update: Node #43, New cost: 1058,
Case 2 Update: Node #65, New cost: 789,
Case 2 Update: Node #188, New cost: 757,
Case 2 Update: Node #245, New cost: 603,
Case 2 Update: Node #1914, New cost: 590,
Case 1 Update: Node #2390, New cost: 458,
Total nodes expanded: 35902, total cost/feasibility look-aheads: 71802
Total flight logic errors caught: 1, total min-days errors caught: 0
Total pruning checks: 101007, total successful: 65107
Best solution: [1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 458

Original problem from client was 06/01/2014 06/07/2014 ATL ORD LAX
Flights ordered by cost: [LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/03/2014, $172]
Flights ordered by date: [ATL->LAX on 06/03/2014, $172, LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151]
Flight check time in hours:minutes:seconds -- 0:6:3.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

We again see no improvement from the $458 route discovered two cases ago. The number of nodes expanded is now at 35,902, but we can still run Balas' algorithm amazingly quickly. We also see our first-ever flight logic error!

As usual, we continue by adding a day. The following input:

```
06/01/2014 06/08/2014 ATL ORD LAX
```

---

[14]If a vector is feasible and has consecutive ones, then there's a good chance it came from a Case 1 update.

has output:

```
Case 1 Update: Node #45, New cost: 1320,
Case 1 Update: Node #81, New cost: 1279,
Case 1 Update: Node #87, New cost: 1058,
Case 2 Update: Node #146, New cost: 789,
Case 2 Update: Node #472, New cost: 757,
Case 2 Update: Node #594, New cost: 603,
Case 2 Update: Node #6098, New cost: 590,
Case 1 Update: Node #7105, New cost: 458,
Total nodes expanded: 127570, total cost/feasibility look-aheads: 255138
Total flight logic errors caught: 1, total min-days errors caught: 0
Total pruning checks: 355251, total successful: 227683
Best solution:
[1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
with cost 458

Original problem from client was 06/01/2014 06/08/2014 ATL ORD LAX
Flights ordered by cost: [LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/03/2014, $172]
Flights ordered by date: [ATL->LAX on 06/03/2014, $172, LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151]
Flight check time in hours:minutes:seconds -- 0:6:47.
DFS search time in hours:minutes:seconds -- 0:0:1.
```

There is no change in the cheapest flight route, but of notice is that Balas' algorithm unfortunately takes about a second to run, with 127,570 nodes expanded.

The following input (we're now approaching double-digit days...):

```
06/01/2014 06/09/2014 ATL ORD LAX
```

has output:

```
Case 1 Update: Node #27, New cost: 1494,
Case 2 Update: Node #37, New cost: 1241,
Case 2 Update: Node #265, New cost: 1067,
Case 1 Update: Node #524, New cost: 1058,
Case 1 Update: Node #769, New cost: 789,
Case 2 Update: Node #1996, New cost: 757,
Case 2 Update: Node #2348, New cost: 603,
Case 2 Update: Node #22578, New cost: 590,
Case 1 Update: Node #26436, New cost: 458,
Total nodes expanded: 460872, total cost/feasibility look-aheads: 921722
Total flight logic errors caught: 20, total min-days errors caught: 0
Total pruning checks: 1280957, total successful: 820106
Best solution: [1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 458

Original problem from client was 06/01/2014 06/09/2014 ATL ORD LAX
Flights ordered by cost: [LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/03/2014, $172]
Flights ordered by date: [ATL->LAX on 06/03/2014, $172, LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/05/2014, $151]
Flight check time in hours:minutes:seconds -- 0:7:36.
DFS search time in hours:minutes:seconds -- 0:0:8.
```

Unfortunately, we get no improvement and Balas' algorithm now takes eight seconds to complete. (Despite our attempts at efficiency, the algorithm *is* exponential.) We now test on a ten-day range with three cities. The input:

```
06/01/2014 06/10/2014 ATL ORD LAX
```

has output:

```
Case 2 Update: Node #30, New cost: 1633,
Case 2 Update: Node #41, New cost: 1380,
Case 2 Update: Node #171, New cost: 1206,
Case 1 Update: Node #408, New cost: 1132,
Case 2 Update: Node #752, New cost: 1052,
Case 1 Update: Node #1616, New cost: 946,
Case 1 Update: Node #5338, New cost: 871,
Case 1 Update: Node #7480, New cost: 728,
Case 1 Update: Node #13964, New cost: 603,
Case 1 Update: Node #37300, New cost: 597,
Case 2 Update: Node #58474, New cost: 580,
Case 2 Update: Node #76216, New cost: 447,
Total nodes expanded: 1472853, total cost/feasibility look-aheads: 2945504
```

```
Total flight logic errors caught: 164, total min-days errors caught: 0
Total pruning checks: 4076472, total successful: 2603784
Best solution:
[1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 447

Original problem from client was 06/01/2014 06/10/2014 ATL ORD LAX
Flights ordered by cost: [LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/10/2014, $140, ATL->LAX on 06/03/2014, $172]
Flights ordered by date: [ATL->LAX on 06/03/2014, $172, LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/10/2014, $140]
Flight check time in hours:minutes:seconds -- 0:8:45.
DFS search time in hours:minutes:seconds -- 0:0:32.
```

Yes! We brought the flight down from \$458 to \$447 with a route that takes us to Los Angeles on June 3, to Chicago on June 4, and then has a layover until it sends us back to Atlanta on June 10. The key here was that extra flight price of \$140. Balas' algorithm goes from needing 8 seconds to needing 32 seconds as we now have to expand almost half a million nodes. We are starting to catch a good number of flight logic errors (164).

Let us add in one more day. The following eleven-day request:

06/01/2014 06/11/2014 ATL ORD LAX

has output:

```
Case 1 Update: Node #34, New cost: 1768,
Case 1 Update: Node #43, New cost: 1515,
Case 1 Update: Node #169, New cost: 1341,
Case 2 Update: Node #398, New cost: 1267,
Case 1 Update: Node #718, New cost: 1187,
Case 2 Update: Node #1624, New cost: 1081,
Case 1 Update: Node #5155, New cost: 1006,
Case 2 Update: Node #7416, New cost: 924,
Case 1 Update: Node #8646, New cost: 892,
Case 2 Update: Node #9058, New cost: 738,
Case 1 Update: Node #22751, New cost: 696,
Case 2 Update: Node #57677, New cost: 603,
Case 2 Update: Node #105804, New cost: 561,
Case 2 Update: Node #213543, New cost: 447,
Case 1 Update: Node #1152833, New cost: 446,
Total nodes expanded: 4609309, total cost/feasibility look-aheads: 9218304
Total flight logic errors caught: 210, total min-days errors caught: 0
Total pruning checks: 12706231, total successful: 8097133
Best solution:
[0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 446

Original problem from client was 06/01/2014 06/11/2014 ATL ORD LAX
Flights ordered by cost: [ORD->LAX on 06/11/2014, $135, ATL->ORD on 06/10/2014, $140, LAX->ATL on 06/03/2014, $171]
Flights ordered by date: [LAX->ATL on 06/03/2014, $171, ATL->ORD on 06/10/2014, $140, ORD->LAX on 06/11/2014, $135]
Flight check time in hours:minutes:seconds -- 0:9:22.
DFS search time in hours:minutes:seconds -- 0:2:23.
```

A *one dollar* improvement! That extra day really did make a difference.

At this point, we will wrap up these experiments. Balas' Algorithm now takes almost two and a half minutes to run, and soon, it will be the limiting factor of our algorithm because the number of nodes to expand will be over ten million for the next iteration. Overall, we are extremely excited with these results. The total runtime of a three-city, eleven day request (so $n(n-1)m = 66$) is still just over half a minute, and this is with using one computer! Twelve cities might take an hour; thirteen cities or more would probably take days.

C.3. Three-City Requests with Minimum Days

Here, we describe some three-city requests *with* minimum days enforced! We subdivide this into two parts. The first will show the output of the following requests:

- 06/01/2014 06/05/2014 ATL ORD LAX 1
- 06/01/2014 06/06/2014 ATL ORD LAX 1
- 06/01/2014 06/07/2014 ATL ORD LAX 1

- 06/01/2014 06/08/2014 ATL ORD LAX 1
- 06/01/2014 06/09/2014 ATL ORD LAX 1

And the second will show the output of requests with more than a day between flights:

- 06/01/2014 06/07/2014 ATL ORD LAX 2
- 06/01/2014 06/08/2014 ATL ORD LAX 2
- 06/01/2014 06/09/2014 ATL ORD LAX 2
- 06/01/2014 06/09/2014 ATL ORD LAX 3

Of note is that our code can detect whether the minimum number of days parameter along with the dates is not sufficient. For instance, the code will immediately abort upon receiving this impossible request:

```
06/01/2014 06/04/2014 ATL ORD LAX 1
```

For brevity, we will further reduce the amount of output shown for each experiment.

C.3.1. THREE-CITY REQUESTS WITH A ONE DAY GAP

A quick note: you can actually tell the request we made because the code will output "Original problem from client...". So in this mini-section (and the next) we'll just refer to them as "Experiment X."

Experiment 1:

```
Total nodes expanded: 2855, total cost/feasibility look-aheads: 5704
Total flight logic errors caught: 4, total min-days errors caught: 66
Total pruning checks: 8159, total successful: 5309
Best solution:
[0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0],
with cost 566

Original problem from client was 06/01/2014 06/05/2014 ATL ORD LAX 1
Flights ordered by date: [ATL->LAX on 06/01/2014, $261, LAX->ORD on 06/03/2014, $154, ORD->ATL on 06/05/2014, $151]
Flight check time in hours:minutes:seconds -- 0:4:23.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

Experiment 2:

```
Total nodes expanded: 10136, total cost/feasibility look-aheads: 20268
Total flight logic errors caught: 2, total min-days errors caught: 148
Total pruning checks: 28703, total successful: 18570
Best solution: [0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 511

Original problem from client was 06/01/2014 06/06/2014 ATL ORD LAX 1
Flights ordered by date: [LAX->ORD on 06/02/2014, $174, ORD->ATL on 06/04/2014, $150, ATL->LAX on 06/06/2014, $187]
Flight check time in hours:minutes:seconds -- 0:5:29.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

Experiment 3:

```
Total nodes expanded: 38613, total cost/feasibility look-aheads: 77218
Total flight logic errors caught: 7, total min-days errors caught: 397
Total pruning checks: 108270, total successful: 69665
Best solution: [0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 482

Original problem from client was 06/01/2014 06/07/2014 ATL ORD LAX 1
Flights ordered by date: [LAX->ORD on 06/03/2014, $154, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/07/2014, $177]
Flight check time in hours:minutes:seconds -- 0:6:7.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

Experiment 4:

```
Total nodes expanded: 125988, total cost/feasibility look-aheads: 251968
Total flight logic errors caught: 7, total min-days errors caught: 1328
Total pruning checks: 350909, total successful: 224929
```

```
Best solution: [0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 482

Original problem from client was 06/01/2014 06/08/2014 ATL ORD LAX 1
Flights ordered by date: [LAX->ORD on 06/03/2014, $154, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/07/2014, $177]
Flight check time in hours:minutes:seconds -- 0:6:41.
DFS search time in hours:minutes:seconds -- 0:0:1.
```

Experiment 5:

```
Total nodes expanded: 442564, total cost/feasibility look-aheads: 885088
Total flight logic errors caught: 31, total min-days errors caught: 5026
Total pruning checks: 1229775, total successful: 787243
Best solution: [1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 468

Original problem from client was 06/01/2014 06/09/2014 ATL ORD LAX 1
Flights ordered by cost: [LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/07/2014, $151, ATL->LAX on 06/09/2014, $182]
Flights ordered by date: [LAX->ORD on 06/04/2014, $135, ORD->ATL on 06/07/2014, $151, ATL->LAX on 06/09/2014, $182]
Flight check time in hours:minutes:seconds -- 0:7:38.
DFS search time in hours:minutes:seconds -- 0:0:8.
```

There are a number of things we can say about these results. The first is that the minimum number of days requirement is respected in all the solutions. Also, the total number of "min-day errors" caught becomes substantial as the date range increases, which is entirely as expected. The number of errors caught is still much smaller than the number of nodes expanded, so it is unlikely that the extra "pruning" here (i.e., pruning in the sense that we discard solutions that violate the minimum days between flight) will result in a noticeable speedup compared to the same request without restriction on minimum number of days. Indeed, looking at the results for this same request but without the day restriction (see Section C.2), both of the nine day requests took eight seconds to go through Balas' algorithm.

Overall, these results show that a user can customize his or her trip to some extent. We chose this customization because it made sense to allow the user to have the option to stay at cities for some minimum length of time. In fact, the lack of this feature this was the number one complaint of our flight system when we explained the basics of it to other Williams students.

C.3.2. THREE-CITY REQUESTS WITH TWO AND THREE DAY GAPS

Here, we show the output of experiments that involve two and three day gaps.

Experiment 1:

```
Total nodes expanded: 38217, total cost/feasibility look-aheads: 76432
Total flight logic errors caught: 2, total min-days errors caught: 457
Total pruning checks: 107555, total successful: 69341
Best solution: [0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 526

Original problem from client was 06/01/2014 06/07/2014 ATL ORD LAX 2
Flights ordered by cost: [ORD->ATL on 06/04/2014, $150, ATL->LAX on 06/07/2014, $172, LAX->ORD on 06/01/2014, $204]
Flights ordered by date: [LAX->ORD on 06/01/2014, $204, ORD->ATL on 06/04/2014, $150, ATL->LAX on 06/07/2014, $172]
Flight check time in hours:minutes:seconds -- 0:6:23.
DFS search time in hours:minutes:seconds -- 0:0:0.
```

Experiment 2:

```
Total nodes expanded: 125267, total cost/feasibility look-aheads: 250532
Total flight logic errors caught: 1, total min-days errors caught: 1411
Total pruning checks: 347973, total successful: 222708
Best solution: [0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 526

Original problem from client was 06/01/2014 06/08/2014 ATL ORD LAX 2
Flights ordered by cost: [ORD->ATL on 06/04/2014, $150, ATL->LAX on 06/07/2014, $172, LAX->ORD on 06/01/2014, $204]
Flights ordered by date: [LAX->ORD on 06/01/2014, $204, ORD->ATL on 06/04/2014, $150, ATL->LAX on 06/07/2014, $172]
Flight check time in hours:minutes:seconds -- 0:6:53.
DFS search time in hours:minutes:seconds -- 0:0:1.
```

Experiment 3:

```
Total nodes expanded: 442122, total cost/feasibility look-aheads: 884242
```

```
Total flight logic errors caught: 2, total min-days errors caught: 5338
Total pruning checks: 1221925, total successful: 779806
Best solution: [0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 507

Original problem from client was 06/01/2014 06/09/2014 ATL ORD LAX 2
Flights ordered by cost: [ORD->ATL on 06/05/2014, $151, LAX->ORD on 06/02/2014, $174, ATL->LAX on 06/09/2014, $182]
Flights ordered by date: [LAX->ORD on 06/02/2014, $174, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/09/2014, $182]
Flight check time in hours:minutes:seconds -- 0:7:53.
DFS search time in hours:minutes:seconds -- 0:0:8.
```

Experiment 4:

```
Total nodes expanded: 442412, total cost/feasibility look-aheads: 884822
Total flight logic errors caught: 1, total min-days errors caught: 30005
Total pruning checks: 1229465, total successful: 787055
Best solution: [0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 537

Original problem from client was 06/01/2014 06/09/2014 ATL ORD LAX 3
Flights ordered by cost: [ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/09/2014, $182, LAX->ORD on 06/01/2014, $204]
Flights ordered by date: [LAX->ORD on 06/01/2014, $204, ORD->ATL on 06/05/2014, $151, ATL->LAX on 06/09/2014, $182]
Flight check time in hours:minutes:seconds -- 0:7:48.
DFS search time in hours:minutes:seconds -- 0:0:9.
```

Again, all the results here are logically correct and obey the minimum days.

C.4. At Least Four Cities

Here, we get into more demanding experiments, designed to test the limits of our computation. First, we run an experiment with the following as input:

```
06/01/2014 06/07/2014 ATL ORD LAX MIA
```

This has $4 \cdot 3 \cdot 7 = 84$ variables, which already exceeds the 66 that were in the three city, eleven day request (which took more than two minutes to finish Balas' algorithm). The output is:

```
Total nodes expanded: 1548120, total cost/feasibility look-aheads: 3096204
Total flight logic errors caught: 24, total min-days errors caught: 0
Total pruning checks: 4598573, total successful: 3050478
Best solution: [1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], with cost 583

Original problem from client was 06/01/2014 06/07/2014 ATL ORD LAX MIA
Flights ordered by cost:
[LAX->ORD on 06/04/2014, $135, MIA->ATL on 06/06/2014, $137,
ORD->MIA on 06/05/2014, $139, ATL->LAX on 06/07/2014, $172]

Flights ordered by date:
[LAX->ORD on 06/04/2014, $135, ORD->MIA on 06/05/2014, $139,
MIA->ATL on 06/06/2014, $137, ATL->LAX on 06/07/2014, $172]

Flight check time in hours:minutes:seconds -- 0:12:21.
DFS search time in hours:minutes:seconds -- 0:0:45.
```

This only took 45 seconds to run, so it is not too bad. Unfortunately, adding in an eighth day resulted in our code getting stuck.

Here is an example of a request with five cities and five days:

```
Total nodes expanded: 261089, total cost/feasibility look-aheads: 522178
Total flight logic errors caught: 0, total min-days errors caught: 0
Total pruning checks: 783093, total successful: 522005
Best solution: [0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
with cost 719

Original problem from client was 06/01/2014 06/05/2014 ATL ORD LAX MIA DFW
```

```
Flights ordered by cost: [DFW->ORD on 06/04/2014, $84, LAX->DFW on 06/03/2014, $118,
ORD->MIA on 06/05/2014, $139, MIA->ATL on 06/01/2014, $177, ATL->LAX on 06/02/2014, $201]

Flights ordered by date: [MIA->ATL on 06/01/2014, $177, ATL->LAX on 06/02/2014, $201,
LAX->DFW on 06/03/2014, $118, DFW->ORD on 06/04/2014, $84, ORD->MIA on 06/05/2014, $139]

Flight check time in hours:minutes:seconds -- 0:14:14.
DFS search time in hours:minutes:seconds -- 0:0:10.
```

So this was not too bad. It took less than 15 minutes to run the whole thing, which had $5 \cdot 4 \cdot 5 = 100$ variables, the most of any experiment in the Appendix. In addition, running an experiment with five cities in six days was also relatively quick (output not shown here). Five cities in seven days would probably be infeasible with just one computer.

```
Flights ordered by cost: [DFW->ORD on 06/04/2014, $84, LAX->DFW on 06/03/2014, $118,
ORD->MIA on 06/05/2014, $139, MIA->ATL on 06/01/2014, $177, ATL->LAX on 06/02/2014, $201]

Flights ordered by date: [MIA->ATL on 06/01/2014, $177, ATL->LAX on 06/02/2014, $201,
LAX->DFW on 06/03/2014, $118, DFW->ORD on 06/04/2014, $84, ORD->MIA on 06/05/2014, $139]
```